


Native RGB + compiled MDD detector

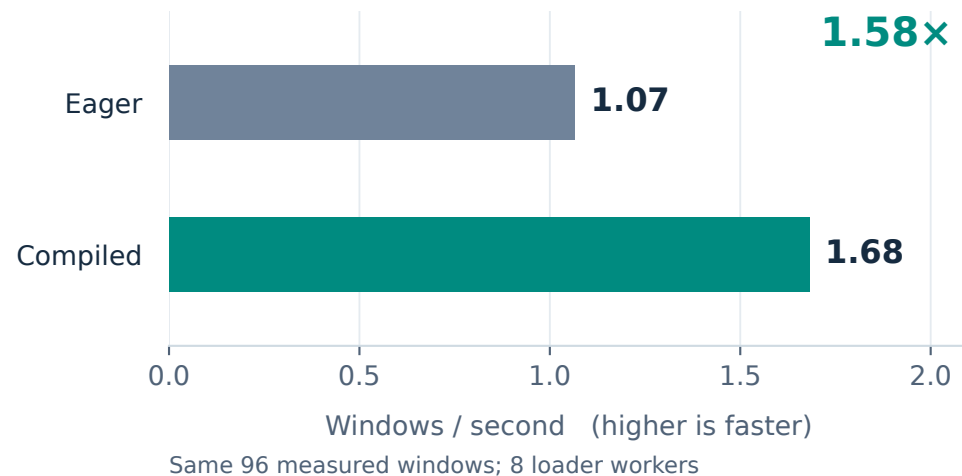
RTX 5060 Ti 16 GB | BF16 | 32 frames at 720p | batch size 1

Same RGB-input model in both conditions; fixed FP32 MDD is inside forward.

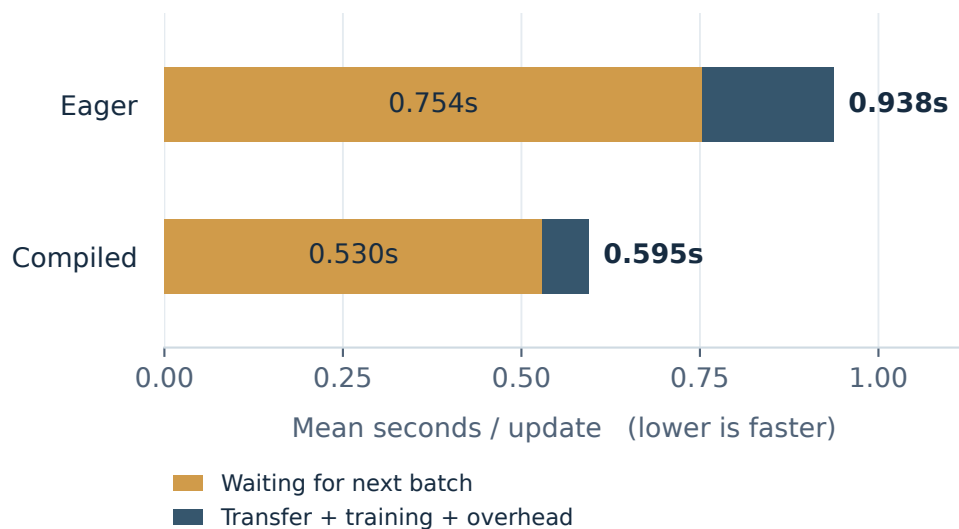
A GPU-resident training



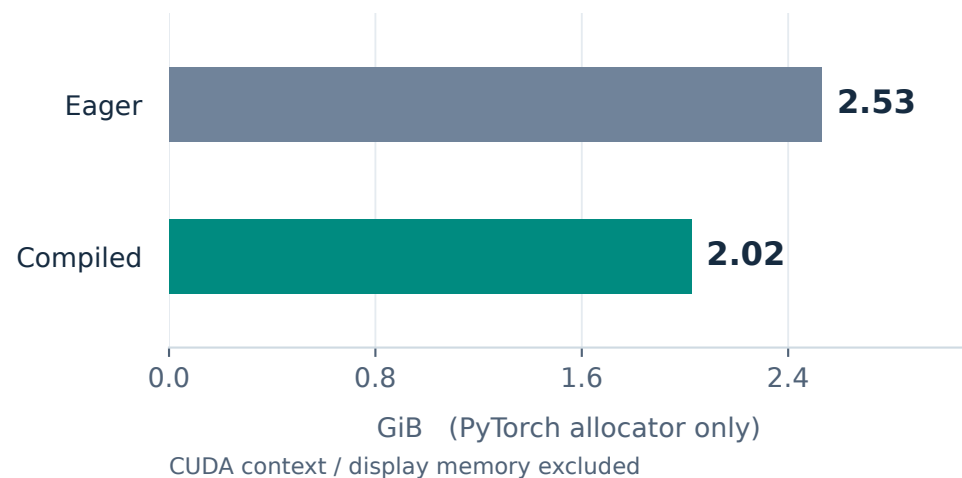
B With CPU reader + transfer



C Wall-clock step with reader



D GPU-resident peak reserved VRAM



One short trial per condition; warmup excluded. Panel A/B scales differ. No accuracy claim.

Source: captured JSON in run-i986-native-rgb-compile-20261009 | PyTorch 2.13.0+cu130 | Inductor default