



Conv2d + query-only | BF16 GPU validation

RTX 5060 Ti 16GB · 1280 × 720 · 32 sampled frames · PyTorch 2.13.0 / CUDA 13.0

A GPU-resident training



B GPU memory



C End-to-end training: same window sequence



Recommended: BS 1, 8 workers. Longer check: 1.33 windows/s over 240 measured updates; peak VRAM 2.40 GiB.

A/B: 4 warm-up + 16 measured updates. C: identical 12 warm-up + 96 measured windows; pin memory, prefetch 1.

VRAM excludes display/context overhead. GPU-resident BS 8: OOM at the 90% allocator cap. Throughput checks, not accuracy results.