



CPU input bottleneck | Conv2d + query-only

1280 × 720 · 32 frames · 8 workers · real frozen train inputs · training remains stopped

A Where CPU preparation time goes



B Isolating the input pipeline



Main targets: CPU luminance/MDD generation and repeated full-clip integrity checks.

A: per-worker service time; workers run concurrently. B: CPU-only; no model, GPU, H2D or pinned-memory timing.

Preverified reader excludes 122.7 s of serial setup (89 clips); it is not an end-to-end speedup claim.

Cached-input case reuses one real tensor to isolate collation/IPC. Other cases use the same 96-window sequence.